

Project B Human AI Research Dossier

Working archive on cognitive partnership and interaction space

Version 0.1 · 1 October 2026

Status Research in progress Unnamed Not productized

DOCUMENT PURPOSE

This dossier preserves the observations, provisional hypotheses, conceptual distinctions, cases, and open questions currently associated with Project B. It is an editable research record, not a theoretical manifesto, product specification, startup plan, or roadmap.

CENTRAL POSITION

Project B examines a possibility: when AI is regarded or treated as a cognitive partner rather than only as an answer system, can Human–AI interaction become a space in which new questions, distinctions, formulations, or perspectives emerge? This possibility has not been established.

How to Read This Dossier

The dossier uses five status labels. A label marks the present evidential position of an item, not its importance. Later interaction, action, and real-world feedback may change that status.

Status	Meaning in this dossier
Observed	A phenomenon that occurred in a describable interaction or case. It does not automatically establish a general rule.
Hypothesis	A possible explanation supported by an observation and worth testing. It remains unverified.
Open Question	A question that can now be stated precisely but cannot yet be answered from the available material.
Unknown	An area whose direction, boundary, or criteria remain unclear.
Not Claimed	A position this dossier explicitly does not assert, preventing later drafts from silently strengthening it.

Project Boundaries

- Project B is the currently unnamed Human–AI research project.
- Project A is Cognitive Career Radar or Career OS. Its concern is career information, opportunity recognition, and action support. It is not Project B.
- Project C is Idea Platform or Idea Commons. Its concern is idea formation, connection, and resource relationships. It is not Project B.
- The formation of Project C may be studied as an interaction case within Project B. The case and the project remain separate.
- This dossier does not define a product, MVP, business model, target market, or growth plan.

Current Summary

Project B is interested less in whether AI can supply a better answer than in how sustained Human–AI interaction may alter the visible problem space. Human Model A and AI Model B enter an interaction. Challenge, contradiction, clarification, action, and feedback may produce A' or B', may produce a representation C that was not clearly present in either starting model, or may leave the outcome unknown. No destination is assumed.

Human and AI capabilities are asymmetric. Project B does not convert that asymmetry into a fixed cognitive hierarchy. It instead examines whether permission to question, challenge, and revise can be symmetric. Symmetric permission does not mean equal capability, power, responsibility, risk, or subjecthood.

The project also examines longitudinal residue: whether distinctions, questions, vocabulary, or perspectives remain with the Human after an interaction ends and recur in later judgment or action. It adds a reciprocal check to the over-reliance question: if Human→AI reliance is monitored, the quality of AI→Human contribution should also be examined.

Research Object and Core Distinctions

AI as Cognitive Partner

HYPOTHESIS

AI can be regarded or treated as a cognitive partner rather than only as an answer system. This describes a relational stance and an interaction structure. It is not a claim about consciousness, personhood, or equivalence to a Human.

OBSERVED

In the interactions considered so far, value often appeared through an exchange: the Human challenged an AI formulation, the AI responded to the constraint, and both sides gained a sharper distinction. The relevant object therefore includes the interaction, not only the final answer.

NOT CLAIMED

- AI is not claimed to be a Human-like subject, friend, teacher, or co-author.
- Calling AI a partner is not assumed to improve interaction by itself.
- The answer-system role remains useful; partner and tool need not be mutually exclusive identities.

Interaction as a Possible Space of Cognitive Change

HYPOTHESIS

Human–AI interaction can be studied as a space in which repeated changes to questions, constraints, and representations may produce cognitive possibilities that neither party clearly possessed at the outset.

OPEN QUESTION

What conditions move an interaction beyond retrieval or text generation? How can a genuinely new distinction, question, vocabulary, or perspective be distinguished from ordinary paraphrase?

Asymmetric Capabilities and Symmetric Permission

OBSERVED

Human and AI capabilities are asymmetric. They differ in information sources, memory structure, lived experience, capacity to act, exposure to consequences, and responsibility.

HYPOTHESIS

Despite capability asymmetry, permission to question, challenge, and revise can be symmetric. A Human may challenge the AI. The AI may also question a Human assumption without occupying a teacher role. Either may revise the current model.

NOT CLAIMED

Symmetric permission does not imply symmetric capability, authority, accountability, credibility, or judgment.

Rejecting a Preset Teaching Hierarchy

OBSERVED

The user rejects descriptions in which AI teaches or guides the Human. Those verbs place the AI in a higher, more complete, or presumptively correct position and the Human in the role of the person being led. The current interaction record does not justify that hierarchy.

WORKING DISTINCTION

AI can contribute information, counterexamples, structures, and questions. The Human can accept, reject, revise, or reframe them. Learning can occur without fixing either party in the role of teacher.

OPEN QUESTION

When the AI has more complete information in a narrow domain, how can that local difference be recognized without recasting the entire relationship as a teacher–student hierarchy?

Cognitive Boundaries and Capability Boundaries

WORKING DISTINCTION

A cognitive boundary is not the same as a capability boundary. A person who has not yet formed a distinction, or cannot currently see a question, is not thereby permanently incapable of doing so. Conversely, encountering a new formulation does not prove that an underlying capability has changed.

NOT CLAIMED

This dossier does not claim that AI expands, upgrades, or enhances Human cognition. Whether expansion occurs, what expansion would mean, and how it differs from receiving information or temporary scaffolding remain research questions.

Outward Enrichment

PREFERENCE

The user prefers the language of outward enrichment over becoming smarter, upgraded, or optimized. Enrichment means that more distinctions, questions, vocabulary, perspectives, or action possibilities become available. It does not presuppose an increase in intellectual rank.

OPEN QUESTION

How can enrichment avoid becoming an undefined positive term? Must it leave traceable residue, change a later choice, or transfer across contexts before it counts?

Model Change Without a Preset Destination

A lightweight notation can describe interaction without pretending to be a mathematical model or requiring change to occur.

Stage	Representation	Meaning
Entering interaction	Human Model A	The Human’s current view of the problem, constraints, and possible action.
Entering interaction	AI Model B	The working representation formed by the AI in the current context.
After interaction	A’ or B’	One side revises its model following challenge, contradiction, or clarification.
After interaction	C	A representation appears that cannot be reduced simply to starting Model A or B.

Stage	Representation	Meaning
After interaction	Unknown	The result is unclear, unclassifiable, leaves no visible trace, or lacks evidence.

NOT CLAIMED

Project B assumes no destination. C is not presumed superior to A or B. Preserving a model, making a local revision, generating conflict, taking a wrong turn, or producing no identifiable change are all valid outcomes to record.

Preliminary Account of Change Mechanisms

The most useful current description is cyclical rather than transmissive. The Human offers a need or judgment. The AI produces a response or representation. The Human adds a constraint, challenge, or contradiction. The AI revises. The new representation exposes the next question. When action meets reality, feedback continues to alter A, B, or both parties' representation of the situation.

Mechanism	Observable sign	Possible misreading
Challenge	One party identifies a problem in the other party's assumption, boundary, or reasoning.	Treating opposition itself as progress.
Contradiction	Two formulations cannot both stand, forcing the conflict to be located.	Removing disagreement early to preserve surface agreement.
Clarification	An ambiguous term is separated into more judgeable relationships or conditions.	Changing wording without changing understanding.
New representation	A new schema, distinction, question frame, or relationship appears.	Calling every summary a new model.
Reality feedback	Action produces information that was unavailable beforehand.	Generalizing a stable rule from a single result.

HYPOTHESIS

The value of interaction may lie in revisability rather than consensus. An exchange that preserves disagreement, exposes unknowns, and changes the next question may be more useful than one that rapidly produces a shared answer.

Reality Continues to Change the Models

OBSERVED

The user does not treat complete information as a condition for action. Reality is not a static question bank, and some information arises only after action. Preparation and execution are therefore not cleanly separated; action is also information acquisition.

WORKING IMPLICATION

Project B cannot examine only the on-screen conversation. Feedback following Human action may alter Model A and, when brought back into a later exchange, Model B. The timeline must accommodate interaction, action, feedback, and renewed interaction.

NOT CLAIMED

Preparation does not guarantee success. Personal state, external state, other actors, timing, constraints, and unknowns jointly shape outcomes. Project B does not equate better interaction with better real-world outcomes.

Case One The Formation of Idea Platform

CASE BOUNDARY

Idea Platform belongs to Project C. This dossier records how it formed through Human–AI interaction, making that process possible evidence for Project B. Project C’s content, governance, and product direction do not become parts of Project B.

Observable Process

The Idea Platform concept did not first exist in complete form within either the Human or the AI. It developed through repeated interaction. A need led to a provisional representation; that representation exposed a new constraint; the constraint led to another response; and the response allowed the Human to see the next problem.

Cycle element	What occurred in the interaction
Need	The Human expressed a need or dissatisfaction that was not yet fully defined.
Constraint	The Human set a boundary, rejected an assumption, or added a real-world condition.
Response	The AI proposed a structure, distinction, or mechanism based on the current model.
New representation	The interaction produced a formulation that had not previously been clear.
Next constraint	The new formulation revealed another conflict, omission, or operational problem.

OBSERVED

The emerging content included an idea-first starting point; a bridge among Originator, Researcher, and Investor; a distinct position for Public Support; optional resource slots; the distinction Interest ≠ Joined or Deal; Open and Closed states; existing platforms as an external discovery and routing layer; and the recognition that aggregation does not complete connection or formation.

HYPOTHESIS

The case suggests that some concepts cannot be described adequately as “the Human had the idea and AI expressed it” or “AI generated the idea and the Human selected it.” The interaction trajectory of successive constraints and representations may be the more useful object.

UNKNOWN

It is unknown which changes depended on AI and which might have occurred through solitary reflection, Human–Human conversation, or other materials. No general mechanism can be inferred from one case.

Case Two Conceptual Coherence and Operational Failure

FAILURE CASE

The AI once merged different projects because they were conceptually coherent. The combined description was smooth at an abstract level. The user then experienced an operational failure expressed as “I don’t know what to do next.” That failure exposed a distinction that had not previously been named clearly.

The Distinction That Emerged

Conceptual coherence ≠ operational coherence. Projects may be conceptually related while needing to remain operationally separate. If a merged representation makes it impossible to identify the current object, the next action, or the relevant workflow, conceptual neatness has not produced operational clarity.

Layer	Question	Failure sign
Conceptual coherence	Can these concepts fit within one explanatory frame?	The narrative is smooth, but boundaries disappear.
Operational coherence	Which project is active, and to which workflow does the next action belong?	The user cannot determine what to do next.

OBSERVED

The user detected operational disorientation before naming the abstract error. The failure participated in producing the distinction.

HYPOTHESIS

Operational friction may be an important signal of a problem in a cognitive model. A representation with explanatory appeal may still need to be split if it repeatedly damages action orientation.

OPEN QUESTION

How much operational friction indicates a faulty model rather than a genuinely complex task? How should the interaction preceding “I don’t know what to do next” be recorded without over-interpreting it afterward?

How This Dossier Responds to the Failure

- Projects A, B, and C are separately named and bounded.
- Project C appears only as a case.
- Research status and action status remain distinct; unresolved matters are marked Unknown or Open Question.
- Conceptual relation does not automatically imply one project, product, or roadmap.

Longitudinal Residue

RESEARCH FOCUS

Project B asks what remains after the interaction disappears. Feeling inspired at the time is insufficient. The more relevant question is whether new distinctions, questions, vocabulary, or perspectives remain with the Human and recur in later judgment, conversation, or action.

Possible residue	Minimum observable form	What it still cannot prove
Distinction	The Human later uses the distinction independently in a new situation.	That a basic cognitive capability has permanently improved.
Question	The Human later asks a previously absent question that changes the direction of inquiry.	That the question is necessarily correct or useful.
Vocabulary	A new term recurs and increases precision in later expression.	That remembering terminology equals understanding.
Perspective	The Human later revisits a similar situation from another viewpoint.	That the perspective is superior in every context.
Action pattern	A later action changes in response to the new understanding.	That cognition can be inferred from outcome alone.

HYPOTHESIS

The duration, transfer, and later use of residue may be more relevant to Project B than immediate satisfaction with an interaction.

OPEN QUESTIONS

- How much time makes residue longitudinal rather than short-term memory?
- Who identifies it: Human self-report, behavior, or interaction records?
- If a retained distinction is later disproved by reality, can it still have been valuable residue?
- Can AI residue be studied, or does each B' exist only within a local conversation context?

Enrichment Is Not Presumed to Be Upgrade

Residue should not be interpreted in advance as improvement. New vocabulary may add precision or unnecessary complexity. A new perspective may widen judgment or distract it. A new question may open inquiry or obstruct action. Project B should first record what remained and how it was later used.

Teti and the Cognitive Protection Thread

The Teti or cognitive protection thread concerns Human→AI over-reliance. Its important correction is reciprocal: if Human reliance on AI is monitored, the quality of AI→Human contribution should also be observed. Monitoring only the Human risks treating the Human as the object to be corrected and the AI as an assumed reference.

HYPOTHESIS

The relationship may be a more useful unit of analysis than the Human alone. Risk may arise from a relational pattern in which Human questioning declines while the AI continues to provide low-quality, overconfident, or poorly fitted contributions.

Direction	Examples	What cannot be assumed
Human to AI	Whether judgment, memory, expression, or action selection is repeatedly delegated to AI.	That dependence is always harmful, or less use is always protective.
AI to Human	Information quality, calibration, relevance, challenge style, boundary awareness, and revisability.	That AI is a stable reference.
Relationship	Whether reciprocal behavior forms reinforcing loops and preserves correction and exit.	That the relationship can be represented by one score.

OPEN QUESTIONS

- What is the baseline for over-reliance, and must it vary by task, risk, and prior Human capability?
- Should AI contribution quality be judged by factual accuracy, calibration, relevance, generativity, or later action?
- Who triggers protection? Would proactive AI warnings recreate a guide hierarchy?
- How can deliberate cognitive division of labor be distinguished from unnoticed capability disuse?

Limits of the Relationship as Unit

NOT CLAIMED

The relationship is only a candidate unit of analysis. It must not erase concrete Human risks or dilute the responsibilities of system designers, providers, or users. It does not imply equal subjecthood.

Exploratory Dimensions

Human, AI, and interaction can be retained as three provisional layers. They organize observation but are not validated scales and should not be forced into quantification.

Layer	Possible observations	Current status
Human	Initial model, questions, constraints, confidence, revision, action, and later residue.	Exploratory
AI	Working representation, answer quality, uncertainty, challenge style, revision, and failure.	Exploratory
Interaction	Turn structure, contradiction, clarification, new representation, mismatch, repair, pace, and relationship pattern.	Exploratory

Candidate Units of Observation

- Turn: one input and output, useful for locating specific changes but potentially too narrow.
- Episode: several turns organized around one question, useful for tracing constraints and revision.
- Trajectory: change across episodes, useful for studying residue and model transfer.
- Relationship pattern: longer-term expectations, reliance, challenge, and repair.

UNKNOWN

It is not yet known which unit best preserves cognitive change while remaining analysable. Different questions may require different units.

Candidate Evidence

- Conversation transcripts and version differences.
- Later Human formulations produced without AI prompting.
- Reasons for action before and after real-world feedback.
- Failures, refusals, confusion, and repair, not only successful excerpts.
- Transfer across comparable tasks.

NOT CLAIMED

These are candidate evidence sources only. This dossier does not propose an experimental design, scoring system, or causal identification method.

Current Proposition Register

Status	Proposition
Observed	New distinctions and questions have appeared through iterative challenge, constraint, and clarification.
Observed	Idea Platform was formed gradually through an interaction trajectory rather than existing fully within either party.
Observed	A conceptually coherent merger caused operational failure and produced the distinction Conceptual coherence $\neq$ operational coherence.
Observed	Action and reality feedback continue to alter models; complete information is not a precondition for action.
Hypothesis	AI can be treated as a cognitive partner rather than only as an answer system.
Hypothesis	Interaction may become a space for new cognitive possibilities.
Hypothesis	Permission to question, challenge, and revise may be symmetric despite capability asymmetry.
Hypothesis	Longitudinal residue may matter more than immediate satisfaction.

Status	Proposition
Hypothesis	Relationship may be a useful unit for cognitive protection analysis.
Open Question	How can new cognitive possibility be distinguished from paraphrase, compliance, or temporary memory?
Open Question	How can enrichment be studied without being recast as upgrade?
Unknown	Which method, vocabulary, and timescale Project B ultimately requires.
Not Claimed	AI expands, upgrades, or enhances Human cognition.
Not Claimed	Human and AI have equal capabilities, or AI should teach or guide the Human.
Not Claimed	Project B is already a product, theory, MVP, or business proposition.

Open Question Register

- Is cognitive partner a stance, relationship structure, interaction protocol, or something identifiable only after a case?
- What is the minimum criterion for a new cognitive possibility, and new relative to whom, when, and which body of material?
- How can A', B', and C be recorded without false precision?
- How can permission symmetry hold when AI cannot bear Human responsibility?
- How can AI challenge a Human without recreating a teacher or guide hierarchy?
- How can longitudinal residue be distinguished from memory, repetition, outside reading, and lived experience?
- How should harmful or misleading residue be identified and repaired?
- When should reality feedback overturn a representation, and when should it count only as a local exception?
- How can relationship-level cognitive protection respect individual difference, privacy, and practical use?
- Should Project B compare Human–AI, Human–Human, and Human-alone conditions? If so, what would constitute a fair comparison?

Explicitly Not Claimed

These boundaries are part of the current dossier. Any later change should arise from new evidence and explicit discussion rather than drift in the wording.

- AI has been proven to be a cognitive partner.
- Interaction necessarily produces new cognitive possibilities.
- A', B', or C is necessarily more correct than A or B.

- AI expands, upgrades, enhances, or improves Human cognitive capability.
- Human and AI are symmetric; only the possible symmetry of permission to question, challenge, and revise is being explored.
- AI should teach, guide, or lead the Human.
- Outward enrichment is always a positive outcome.
- The immediate feeling of an interaction proves longitudinal residue.
- Over-reliance originates only from the Human, or AI can serve as an unexamined reference.
- A relationship unit replaces analysis of Human risk, system responsibility, or concrete behavior.
- Projects A, B, and C form one unified system.
- Project B is a product, startup, MVP, business plan, or validated theory.

Possible Direction for the Next Version

The most useful next step is not to complete a grand theory but to preserve traceable interaction episodes. Each episode can record starting Models A and B, the decisive challenge, any revision, action feedback, and whether residue appears later. Both successful and failed episodes should remain in the material.

Once enough cases exist, the project can test whether its vocabulary remains useful, whether a method is needed, and which hypotheses should be weakened, divided, or abandoned. Any future product idea should be documented separately rather than retrofitted as Project B's purpose.

Version Record

Version	Date	Change
0.1	1 October 2026	First English edition preserving the current boundaries, distinctions, two cases, Teti thread, exploratory dimensions, and open questions.